



HBase

Shell commands

Data Ingestion

Integration

Type hbase in shell:

Commands:

Some commands take arguments. Pass no args or -h for usage.

shell	Run the HBase shell
hbck	Run the hbase 'fsck' tool
snapshot	Create a new snapshot of a table
snapshotinfo	Tool for dumping snapshot information
wal	Write-ahead-log analyzer
hfile	Store file analyzer
zkcli	Run the ZooKeeper shell
upgrade	Upgrade hbase
master	Run an HBase HMaster node
regionserver	Run an HBase HRegionServer node
zookeeper	Run a Zookeeper server
rest	Run an HBase REST server
thrift	Run the HBase Thrift server
thrift2	Run the HBase Thrift2 server
clean	Run the HBase clean up script
classpath	Dump hbase CLASSPATH
mapredcp	Dump CLASSPATH entries required by mapreduce
pe	Run PerformanceEvaluation
ltt	Run LoadTestTool
version	Print the version
CLASSNAME	Run the class named CLASSNAME

You see many options!

For developer, it is shell!

Lets deep dive in to shell.

Type hbase shell

```
pankaj@ip-172-31-4-182:~$ hbase shell
17/01/06 04:16:56 INFO Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.8.0, rUnknown, Tue Jul 12 16:02:46 PDT 2016
```

```
hbase(main):001:0> █
```

Type hbase help

COMMAND GROUPS:

Group name: general

Commands: status, table_help, version, whoami

Group name: ddl

Commands: alter, alter_async, alter_status, create, describe, disable, disable_all, drop, drop_all, enable, enable_all, exists, get_table, is_disabled, is_enabled, list, locate_region, show_filters

Group name: namespace

Commands: alter_namespace, create_namespace, describe_namespace, drop_namespace, list_namespace, list_namespace_tables

Group name: dml

Commands: append, count, delete, deleteall, get, get_counter, get_splits, incr, put, scan, truncate, truncate_preserve

Group name: tools

Commands: assign, balance_switch, balancer, balancer_enabled, catalogjanitor_enabled, catalogjanitor_run, catalogjanitor_switch, close_region, compact, compact_mob, compact_rs, flush, major_compact, major_compact_mob, merge_region, move, normalize, normalizer_enabled, normalizer_switch, split, trace, unassign, wal_roll, zk_dump

Group name: replication

Commands: add_peer, append_peer_tableCFs, disable_peer, disable_table_replication, enable_peer, enable_table_replication, list_peers, list_replicated_tables, remove_peer, remove_peer_tableCFs, set_peer_tableCFs, show_peer_tableCFs

hbase status

```
hbase(main):002:0> hbase status
1 active master, 0 backup masters, 3 servers, 0 dead, 21.6667 average load
```

hbase version

```
hbase(main):003:0> hbase version
1.2.0-cdh5.8.0, rUnknown, Tue Jul 12 16:02:46 PDT 2016
```

hbase table_help

```
hbase(main):004:0> hbase table_help
Help for table-reference commands.
```

You can either create a table via 'create' and then manipulate the table via commands like 'put', 'get', etc. See the standard help information for how to use each of these commands.

However, as of 0.96, you can also get a reference to a table, on which you can invoke commands. For instance, you can get create a table and keep around a reference to it via:

hbase whoami

```
hbase(main):007:0> hbase whoami
pankaj (auth:SIMPLE)
groups: labusers
```

Hbase DML and DDL command line operations

Following Operations will performed:

1. Put- Data Ingestion
2. Data Retrieval:
 1. count- number of records
 2. get – Retrieve a row
 3. scan – Read entire table
 4. delete - delete a particular row
 5. put – update table
3. Drop table:
 1. disable – Notifying region servers that table will be removed
 2. drop – table will be dropped by master
4. Refer to **commandLineHbase.pdf** file for lab activities.

Sample casestudy

Consider a telecom sector is having 100 million customers. Now the challenge is how to store all possible information of a given customer.

Solution is on NoSQL:

Let us create a table with name, **telecom_username** with following column families. Fill in your name in place of username:

1. Demographics – Database table
2. Plansactive - JSON File
3. Latest3mails – TSV File
4. Paymentsinfo - Database table

*** Data created is only for practice.

Create table on HBase

hbase shell

```
create 'telco_username',  
{NAME=>'Demographics'}, {NAME=>'Plansactive'}, {NAME=>'Latest3mails'}, {NAME=>'Paymentsinfo'}
```

```
list 'telco.*'
```

```
hbase(main):001:0> create 'telco_instructor', {NAME=>'Demographics'}, {NAME=>'Plansactive'}, {NAME=>'Latest3mails'}, {NAME=>'Paymentsinfo'}  
0 row(s) in 2.4130 seconds  
  
=> Hbase::Table - telco_instructor  
hbase(main):002:0> list 'telco.*'  
TABLE
```

Ingestion of Structured Data using Sqoop

Before ingesting data, create tables on MySQL.

```
--login to mysql:
```

```
mysql -h 54.149.41.179 -u username -p  
use database;
```

```
--create tables:
```

```
create table Demographics (custid varchar(10),age int, gender varchar(1));
```

```
create table paymentsinfo (custid varchar(10),ontime1 int,delayed1  
int,total1 int);
```

```
mysql> create table Demographics (custid varchar(10),age int, gender varchar(1));  
Query OK, 0 rows affected (0.01 sec)
```

```
mysql>  
mysql> create table paymentsinfo (custid varchar(10),ontime1 int,delayed1 int,total1 int);  
Query OK, 0 rows affected (0.01 sec)
```


Ingestion of Structured Data using Sqoop contd..

--Insert into demographics:

```
insert into Demographics (custid,age,gender) values('9065267058',46,'m');
insert into Demographics (custid,age,gender) values('9314449592',20,'f');
insert into Demographics (custid,age,gender) values('9514131750',21,'f');
insert into Demographics (custid,age,gender) values('9888944896',25,'f');
insert into Demographics (custid,age,gender) values('9566797883',36,'m');
insert into Demographics (custid,age,gender) values('9569949682',47,'f');
insert into Demographics (custid,age,gender) values('9002620762',28,'f');
insert into Demographics (custid,age,gender) values('9266241062',49,'m');
insert into Demographics (custid,age,gender) values('9058015369',20,'m');
insert into Demographics (custid,age,gender) values('9130634376',23,'m');
insert into Demographics (custid,age,gender) values('9725923789',21,'m');
```

Ingestion of Structured Data using Sqoop contd..

--Insert into paymentsinfo:

```
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9065267058',16,43,59);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9314449592',24,47,71);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9514131750',32,19,51);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9888944896',11,37,48);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9566797883',38,38,76);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9569949682',18,41,59);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9002620762',30,20,50);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9266241062',23,36,59);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9058015369',36,26,62);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9130634376',19,43,62);
Insert into paymentsinfo(custid ,ontime1 ,delayed1 ,total1 ) values ( '9725923789',16,38,54);
```

MYSQL <--> SQOOP <--> HBASE

Enter dbname and username, to import Demographics table using sqoop

```
sqoop import \  
--connect jdbc:mysql://54.149.41.179/dbname --username username -P \  
--table Demographics \  
--columns "custid,age,gender" \  
--hbase-table telco_username \  
--column-family Demographics \  
--hbase-row-key custid -m 1
```

Run the following command in HBase:

```
get 'telco_username', 9065267058  
hbase(main):002:0> get 'telco_instructor', 9065267058  
COLUMN                                CELL  
  Demographics:age                    timestamp=1483978028908, value=46  
  Demographics:gender                 timestamp=1483978028908, value=m  
2 row(s) in 0.0440 seconds
```

Enter dbname and username, to import paymentsinfo table using sqoop

```
sqoop import \  
--connect jdbc:mysql://54.149.41.179/dbname --username username -P \  
--table paymentsinfo \  
--columns "custid,ontime1,delayed1,total1" \  
--hbase-table telco_username \  
--column-family Paymentsinfo \  
--hbase-row-key custid -m 1
```

Run the following command in HBase:

get 'telco_username', 9725923789

```
hbase(main):003:0> get 'telco_instructor', 9725923789  
COLUMN                                CELL  
Demographics:age                      timestamp=1483978028908, value=21  
Demographics:gender                   timestamp=1483978028908, value=m  
Paymentsinfo:delayed1                 timestamp=1483981728892, value=38  
Paymentsinfo:ontime1                  timestamp=1483981728892, value=16  
Paymentsinfo:total1                   timestamp=1483981728892, value=54  
5 row(s) in 0.0080 seconds
```

Ingestion of Raw CSV file using JAVA API

Using winscp, copy email.csv file to cluster

Create a directory on HDFS:

```
hadoop fs -mkdir hbasepractice
```

```
hadoop fs -put email.csv hbasepractice
```

```
hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator=, -
Dimporttsv.columns=HBASE_ROW_KEY,Latest3mails:mail3,Latest3mails:mail2,Latest3mails:mail
1 telco_username hbasepractice/email.csv
```

Run the following command in HBase:

```
get 'telco_username', 9888944896
```

```
hbase(main):004:0> get 'telco_instructor', 9888944896
COLUMN                                CELL
Demographics:age                      timestamp=1483978028908, value=25
Demographics:gender                   timestamp=1483978028908, value=f
Latest3mails:mail1                    timestamp=1483982058013, value=
Latest3mails:mail2                    timestamp=1483982058013, value=
Latest3mails:mail3                    timestamp=1483982058013, value=xxxx xxx xx
Paymentsinfo:delayed1                 timestamp=1483981728892, value=37
Paymentsinfo:ontime1                  timestamp=1483981728892, value=11
Paymentsinfo:total1                   timestamp=1483981728892, value=48
8 row(s) in 0.0120 seconds
```

Ingestion of semistructured data using Apache Pig

Using WinSCP, copy the plans.json file to cluster.

```
hadoop fs -put plans.json hbasepractice
```

Login to Pig

```
Plansactive = LOAD 'hbasepractice/plans.json' USING
JsonLoader('custid:chararray, messageplan:chararray,
            dataplan:chararray,      callplan:chararray,
            comboplan:chararray');
```

```
STORE Plansactive INTO 'hbase://telco_username' USING
org.apache.pig.backend.hadoop.hbase.HBaseStorage(
' Plansactive:messageplan
Plansactive:dataplan
Plansactive:callplan
Plansactive:comboplan'
);
```

Run the following command in HBase:

```
get 'telco_username', 9065267058
```

```
hbase(main):002:0> get 'telco_instructor', 9065267058
COLUMN                                CELL
Demographics:age                      timestamp=1483978028908, value=46
Demographics:gender                   timestamp=1483978028908, value=m
Latest3mails:mail1                   timestamp=1483982058013, value=
Latest3mails:mail2                   timestamp=1483982058013, value=xxxx xxx xx
Latest3mails:mail3                   timestamp=1483982058013, value=xxxx xxx xx
Plansactive:callplan                 timestamp=1483985217693, value=yes
Plansactive:dataplan                 timestamp=1483985217693, value=
Plansactive:key                       timestamp=1483985217693, value=yes
Plansactive:messageplan              timestamp=1483985217693, value=
9 row(s) in 0.0550 seconds
```

Access table with Hive

Now on this table, access can be provided to specific information to different departments.

All helpdesk should know is, your demographics information. Lets enable it:

```
create database hbase_hive_username;
use hbase_hive_username;

create external table if not exists helpdesk(custid string, age int,
gender string)
stored by 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES("hbase.columns.mapping" = ":key, Demographics:age,
Demographics:gender")
TBLPROPERTIES("hbase.table.name"="telco_username");

select count(*) from helpdesk where age>30;

select * from helpdesk where custid= 9065267058;
```

Access table with Hive – Marketing Department

```
create external table if not exists marketing (custid int, age int,  
gender string,  
messageplan string, dataplan string, callplan string, comboplan string)  
stored by 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
  WITH SERDEPROPERTIES("hbase.columns.mapping" = ":key,  
Demographics:age,  
Demographics:gender,  
Plansactive:messageplan,  
Plansactive:dataplan,  
Plansactive:callplan,  
Plansactive:comboplan ")  
TBLPROPERTIES("hbase.table.name"="telco_username");  
  
Select * from marketing limit 10;  
-- Number of customer who hold messageplan  
  
Select count(*) from marketing where messageplan is not null;
```


Conclusion

In this lab, you have learnt the following:

1. HBase shell commands
2. Sqoop <--> HBase integration
3. Pig <--> HBase integration
4. Hive <--> HBase integration
5. Raw file load to HBase